

# Regression on Boaston dataset

- SHWETA KANHERE BANAIT
  - GitHub :<https://github.com/shwetawin/PythonProject> (<https://github.com/shwetawin/PythonProject>)
  - Linkdin: <https://www.linkedin.com/in/shweta-kanhere-07810213/> (<https://www.linkedin.com/in/shweta-kanhere-07810213/>)
  - Email:[shweta.kanhere@gmail.com](mailto:shweta.kanhere@gmail.com) (<mailto:shweta.kanhere@gmail.com>)

## Types of Regression

1. Linear Regression
2. Ridge Regression
3. Lasso Regression
4. Elastic-Net Regression

In [1]:

```
1 import pandas as pd
2 import matplotlib.pyplot as plt
3 import numpy as np
4 import seaborn as sns
5 from sklearn.datasets import load_boston
```

In [2]:

```
1 boston= load_boston()
```

C:\Users\Shweta Kanhere\anaconda3\lib\site-packages\sklearn\utils\deprecation.py:87: FutureWarning: Function load\_boston is deprecated; `load\_boston` is deprecated in 1.0 and will be removed in 1.2.

The Boston housing prices dataset has an ethical problem. You can refer to the documentation of this function for further details.

The scikit-learn maintainers therefore strongly discourage the use of this dataset unless the purpose of the code is to study and educate about ethical issues in data science and machine learning.

In this special case, you can fetch the dataset from the original source::

```
import pandas as pd
import numpy as np
```

```
data_url = "http://lib.stat.cmu.edu/datasets/boston"
raw_df = pd.read_csv(data_url, sep="\s+", skiprows=22, header=None)
data = np.hstack([raw_df.values[::2, :], raw_df.values[1::2, :2]])
target = raw_df.values[1::2, 2]
```

Alternative datasets include the California housing dataset (i.e. :func:`~sklearn.datasets.fetch\_california\_housing`) and the Ames housing dataset. You can load the datasets as follows::

```
from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing()
```

for the California housing dataset and::

```
from sklearn.datasets import fetch_openml
housing = fetch_openml(name="house_prices", as_frame=True)
```

for the Ames housing dataset.

```
warnings.warn(msg, category=FutureWarning)
```

In [3]:

```
1 # to get keys of the data set
2 boston.keys()
```

Out[3]:

```
dict_keys(['data', 'target', 'feature_names', 'DESCR', 'filename', 'data_module'])
```

In [4]:

```
1 # data set description
2 print(boston.DESCR)
```

```
.. _boston_dataset:
```

Boston house prices dataset

-----

**\*\*Data Set Characteristics:\*\***

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B  $1000(B_k - 0.63)^2$  where  $B_k$  is the proportion of black people by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/> (<http://archive.ics.uci.edu/ml/machine-learning-databases/housing/>)

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [5]:

```
1 # to get input feature data
2 print(boston.data)
```

```
[[6.3200e-03 1.8000e+01 2.3100e+00 ... 1.5300e+01 3.9690e+02 4.9800e+00]
 [2.7310e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9690e+02 9.1400e+00]
 [2.7290e-02 0.0000e+00 7.0700e+00 ... 1.7800e+01 3.9283e+02 4.0300e+00]
 ...
 [6.0760e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 5.6400e+00]
 [1.0959e-01 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9345e+02 6.4800e+00]
 [4.7410e-02 0.0000e+00 1.1930e+01 ... 2.1000e+01 3.9690e+02 7.8800e+00]]
```

In [6]:

```
1 # to get output feature means target (here output feature is cost of house which is target)
2 print(boston.target)
```

```
[24. 21.6 34.7 33.4 36.2 28.7 22.9 27.1 16.5 18.9 15. 18.9 21.7 20.4
18.2 19.9 23.1 17.5 20.2 18.2 13.6 19.6 15.2 14.5 15.6 13.9 16.6 14.8
18.4 21. 12.7 14.5 13.2 13.1 13.5 18.9 20. 21. 24.7 30.8 34.9 26.6
25.3 24.7 21.2 19.3 20. 16.6 14.4 19.4 19.7 20.5 25. 23.4 18.9 35.4
24.7 31.6 23.3 19.6 18.7 16. 22.2 25. 33. 23.5 19.4 22. 17.4 20.9
24.2 21.7 22.8 23.4 24.1 21.4 20. 20.8 21.2 20.3 28. 23.9 24.8 22.9
23.9 26.6 22.5 22.2 23.6 28.7 22.6 22. 22.9 25. 20.6 28.4 21.4 38.7
43.8 33.2 27.5 26.5 18.6 19.3 20.1 19.5 19.5 20.4 19.8 19.4 21.7 22.8
18.8 18.7 18.5 18.3 21.2 19.2 20.4 19.3 22. 20.3 20.5 17.3 18.8 21.4
15.7 16.2 18. 14.3 19.2 19.6 23. 18.4 15.6 18.1 17.4 17.1 13.3 17.8
14. 14.4 13.4 15.6 11.8 13.8 15.6 14.6 17.8 15.4 21.5 19.6 15.3 19.4
17. 15.6 13.1 41.3 24.3 23.3 27. 50. 50. 50. 22.7 25. 50. 23.8
23.8 22.3 17.4 19.1 23.1 23.6 22.6 29.4 23.2 24.6 29.9 37.2 39.8 36.2
37.9 32.5 26.4 29.6 50. 32. 29.8 34.9 37. 30.5 36.4 31.1 29.1 50.
33.3 30.3 34.6 34.9 32.9 24.1 42.3 48.5 50. 22.6 24.4 22.5 24.4 20.
21.7 19.3 22.4 28.1 23.7 25. 23.3 28.7 21.5 23. 26.7 21.7 27.5 30.1
44.8 50. 37.6 31.6 46.7 31.5 24.3 31.7 41.7 48.3 29. 24. 25.1 31.5
23.7 23.3 22. 20.1 22.2 23.7 17.6 18.5 24.3 20.5 24.5 26.2 24.4 24.8
29.6 42.8 21.9 20.9 44. 50. 36. 30.1 33.8 43.1 48.8 31. 36.5 22.8
30.7 50. 43.5 20.7 21.1 25.2 24.4 35.2 32.4 32. 33.2 33.1 29.1 35.1
45.4 35.4 46. 50. 32.2 22. 20.1 23.2 22.3 24.8 28.5 37.3 27.9 23.9
21.7 28.6 27.1 20.3 22.5 29. 24.8 22. 26.4 33.1 36.1 28.4 33.4 28.2
22.8 20.3 16.1 22.1 19.4 21.6 23.8 16.2 17.8 19.8 23.1 21. 23.8 23.1
20.4 18.5 25. 24.6 23. 22.2 19.3 22.6 19.8 17.1 19.4 22.2 20.7 21.1
19.5 18.5 20.6 19. 18.7 32.7 16.5 23.9 31.2 17.5 17.2 23.1 24.5 26.6
22.9 24.1 18.6 30.1 18.2 20.6 17.8 21.7 22.7 22.6 25. 19.9 20.8 16.8
21.9 27.5 21.9 23.1 50. 50. 50. 50. 50. 13.8 13.8 15. 13.9 13.3
13.1 10.2 10.4 10.9 11.3 12.3 8.8 7.2 10.5 7.4 10.2 11.5 15.1 23.2
9.7 13.8 12.7 13.1 12.5 8.5 5. 6.3 5.6 7.2 12.1 8.3 8.5 5.
11.9 27.9 17.2 27.5 15. 17.2 17.9 16.3 7. 7.2 7.5 10.4 8.8 8.4
16.7 14.2 20.8 13.4 11.7 8.3 10.2 10.9 11. 9.5 14.5 14.1 16.1 14.3
11.7 13.4 9.6 8.7 8.4 12.8 10.5 17.1 18.4 15.4 10.8 11.8 14.9 12.6
14.1 13. 13.4 15.2 16.1 17.8 14.9 14.1 12.7 13.5 14.9 20. 16.4 17.7
19.5 20.2 21.4 19.9 19. 19.1 19.1 20.1 19.9 19.6 23.2 29.8 13.8 13.3
16.7 12. 14.6 21.4 23. 23.7 25. 21.8 20.6 21.2 19.1 20.6 15.2 7.
8.1 13.6 20.1 21.8 24.5 23.1 19.7 18.3 21.2 17.5 16.8 22.4 20.6 23.9
22. 11.9]
```

In [7]:

```
1 # to get feature name
2 print(boston.feature_names)
```

```
['CRIM' 'ZN' 'INDUS' 'CHAS' 'NOX' 'RM' 'AGE' 'DIS' 'RAD' 'TAX' 'PTRATIO'
'B' 'LSTAT']
```

In [8]:

```
1 print(boston.filename)
```

boston\_house\_prices.csv

In [9]:

```
1 print(boston.data_module)
```

sklearn.datasets.data

In [10]:

```
1 ###Preapre the dataframe
2
```

In [11]:

```
1 dataset=pd.DataFrame(boston.data,columns=boston.feature_names)
2 dataset.head()
```

Out[11]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	1

In [12]:

```
1 # to get the information about dataset
2 dataset.info()
```

&lt;class 'pandas.core.frame.DataFrame'&gt;

RangeIndex: 506 entries, 0 to 505

Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	CRIM	506 non-null	float64
1	ZN	506 non-null	float64
2	INDUS	506 non-null	float64
3	CHAS	506 non-null	float64
4	NOX	506 non-null	float64
5	RM	506 non-null	float64
6	AGE	506 non-null	float64
7	DIS	506 non-null	float64
8	RAD	506 non-null	float64
9	TAX	506 non-null	float64
10	PTRATIO	506 non-null	float64
11	B	506 non-null	float64
12	LSTAT	506 non-null	float64

dtypes: float64(13)

memory usage: 51.5 KB

In [13]:

```
1 dataset.shape
```

Out[13]:

(506, 13)

In [14]:

```
1 # Add target variable Price
2 dataset["Price"]=boston.target
```

In [15]:

```
1 dataset.head()
```

Out[15]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	1
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	1

# EDA

In [16]:

```
1 dataset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  --
 0   CRIM        506 non-null    float64
 1   ZN          506 non-null    float64
 2   INDUS       506 non-null    float64
 3   CHAS        506 non-null    float64
 4   NOX         506 non-null    float64
 5   RM          506 non-null    float64
 6   AGE         506 non-null    float64
 7   DIS         506 non-null    float64
 8   RAD         506 non-null    float64
 9   TAX         506 non-null    float64
10  PTRATIO     506 non-null    float64
11  B           506 non-null    float64
12  LSTAT       506 non-null    float64
13  Price       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB
```

In [17]:

```
1 dataset.shape
```

Out[17]:

```
(506, 14)
```



In [18]:

```
1 dataset.describe().T
```

Out[18]:

	count	mean	std	min	25%	50%	75%	max
<b>CRIM</b>	506.0	3.613524	8.601545	0.00632	0.082045	0.25651	3.677083	88.976
<b>ZN</b>	506.0	11.363636	23.322453	0.00000	0.000000	0.00000	12.500000	100.000
<b>INDUS</b>	506.0	11.136779	6.860353	0.46000	5.190000	9.69000	18.100000	27.740
<b>CHAS</b>	506.0	0.069170	0.253994	0.00000	0.000000	0.00000	0.000000	1.000
<b>NOX</b>	506.0	0.554695	0.115878	0.38500	0.449000	0.53800	0.624000	0.871
<b>RM</b>	506.0	6.284634	0.702617	3.56100	5.885500	6.20850	6.623500	8.780
<b>AGE</b>	506.0	68.574901	28.148861	2.90000	45.025000	77.50000	94.075000	100.000
<b>DIS</b>	506.0	3.795043	2.105710	1.12960	2.100175	3.20745	5.188425	12.120
<b>RAD</b>	506.0	9.549407	8.707259	1.00000	4.000000	5.00000	24.000000	24.000
<b>TAX</b>	506.0	408.237154	168.537116	187.00000	279.000000	330.00000	666.000000	711.000
<b>PTRATIO</b>	506.0	18.455534	2.164946	12.60000	17.400000	19.05000	20.200000	22.000
<b>B</b>	506.0	356.674032	91.294864	0.32000	375.377500	391.44000	396.225000	396.900
<b>LSTAT</b>	506.0	12.653063	7.141062	1.73000	6.950000	11.36000	16.955000	37.970
<b>Price</b>	506.0	22.532806	9.197104	5.00000	17.025000	21.20000	25.000000	50.000

- see ( 25 -50- 75 )%
- Chas has no outliers
- INDUS has few outliers
- Price has minute outliers

In [19]:

```
1 # check missing values
2 dataset.isnull().sum()
```

Out[19]:

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
Price     0
dtype: int64
```

In [20]:

```
1 ### no missing values
```

- Information
- In Linear regression the most important thing is the relation ship between dependent and Independent features

In [21]:

```
1 # EDA
2 dataset.corr()
```

Out[21]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	-0.379670
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	0.664408
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	-0.708027
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	-0.099176
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	-0.769230
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	0.205246
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	-0.747881
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	1.000000
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	-0.494588
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	-0.534432
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	-0.232471
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	0.291512
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	-0.496996
Price	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	0.249929

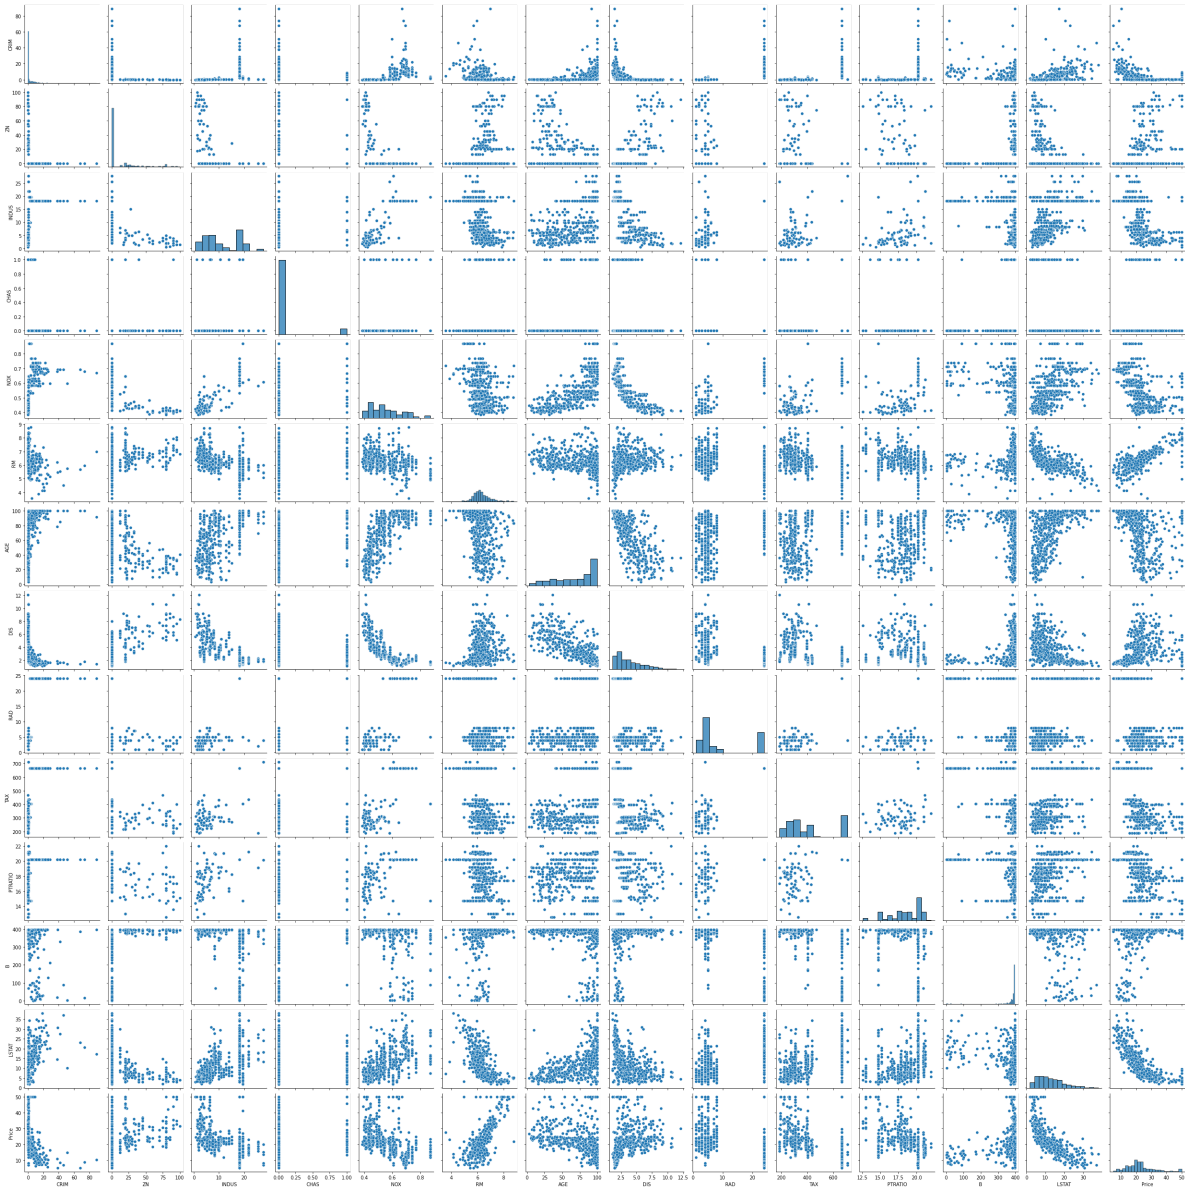
- Relationship betn CRIM and ZN is -0.2004
- CRIM increase Price decreases
- Tax and RAD has good coorelation
- model is good only if relation between dependent and indepent variable is good

In [22]:

```
1 sns.pairplot(dataset)
```

Out[22]:

<seaborn.axisgrid.PairGrid at 0x2069339b1c0>




In [23]:

```
1 ## can see this above in better way
2 sns.heatmap(dataset.corr())
```

Out[23]:

&lt;AxesSubplot:&gt;




In [24]:

```
1 ## size of graph will increase by figure size  
2 sns.set(rc={'figure.figsize':(15,15)})  
3 sns.heatmap(dataset.corr())
```

Out[24]:

&lt;AxesSubplot:&gt;





In [25]:


```

1  ## chek the % by annot= True
2  sns.set(rc={'figure.figsize':(15,15)})
3  sns.heatmap(dataset.corr(),annot= True)

```

Out[25]:

&lt;AxesSubplot:&gt;






In [26]:

```
1 # to see some relationship betn crim and Price
2 # use scatter plot
3 plt.scatter(dataset["CRIM"],dataset["Price"])
4 plt.title(" Analysis of House Price WRT CRIME")
5 plt.xlabel("Crime rate of the city")
6 plt.ylabel("Housing Price of the city")
7
```

Out[26]:

Text(0, 0.5, 'Housing Price of the city')




Observation : crime is high then price is low

In [27]:

```
1 plt.scatter(dataset["RM"],dataset["Price"])
2 plt.title(" Analysis of House Price WRT rooms per dwelling")
3 plt.xlabel("Caverage number of rooms per dwelling")
4 plt.ylabel("Housing Price of the city")
```

Out[27]:

Text(0, 0.5, 'Housing Price of the city')



In [94]:

```

1  ### Visualising all the features with respect to price
2  for feature in [feature for feature in dataset.columns if feature not in ['Price']]:
3      sns.set(rc={'figure.figsize':(8,8)})
4      plt.scatter(x=dataset[feature], y=dataset['Price'])
5      plt.xlabel(feature)
6      plt.ylabel("Price")
7      plt.title("{} Vs Price".format(feature))
8      plt.show();

```




- good correlaton between them
  - RM is good feature to find output
  - model will create good feet line
- 
- The REGPLOT----- Plot data in a linear regression model fit. There are a number of mutually exclusive options for estimating the regression model

In [28]:

```
1 sns.set(rc={'figure.figsize':(10,10)}) # set the figure size
2 sns.regplot(data=dataset,x="RM",y="Price")
```

Out[28]:

&lt;AxesSubplot:xlabel='RM', ylabel='Price'&gt;




- shaded region shows region lasso
- with change in Lamda we get best fit line
- concentration of points are more then movement is less and concentration of points are less then movement is high

In [29]:

```
1 #LSTAT - % lower status of the population
2 sns.regplot(data=dataset,x="LSTAT",y="Price")
```

Out[29]:

<AxesSubplot:xlabel='LSTAT', ylabel='Price'>




In [30]:

```
1 sns.set(rc={'figure.figsize':(6,6)}) # set the figure size
2 sns.regplot(data=dataset,x="CRIM",y="Price") #----- negative correlation
```


Out[30]:

&lt;AxesSubplot:xlabel='CRIM', ylabel='Price'&gt;



In [95]:

```
1 ##### Regression plot for all other features wrt Price
2 for feature in [feature for feature in dataset.columns if feature not in ['Price']]:
3     sns.set(rc={'figure.figsize':(8,8)})
4     sns.regplot(x=dataset[feature], y=dataset['Price'])
5     plt.xlabel(feature)
6     plt.ylabel("Price")
7     plt.title("{} Vs Price".format(feature))
8     plt.show();
```



In [31]:


```
1 sns.boxplot(dataset['Price'])
```

C:\Users\Shweta Kanhere\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[31]:

<AxesSubplot:xlabel='Price'>



Less outliers in Price



In [32]:


```
1 sns.boxplot(dataset['CRIM'])
```

C:\Users\Shweta Kanhere\anaconda3\lib\site-packages\seaborn\\_decorators.py:36: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

Out[32]:

<AxesSubplot:xlabel='CRIM'>



In [33]:

```
1 ### More outliers with crime ant this impact Price
```

In [96]:

```

1 # Box plot for all the features
2 for feature in dataset.columns:
3     sns.set(rc={'figure.figsize':(8,8)})
4     sns.boxplot(dataset[feature])
5     plt.xlabel(feature)
6     plt.title("{}".format(feature))
7     plt.show();

```

version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

```
warnings.warn(
```

## Training the Model

In [34]:

```
1 dataset.head()
```

Out[34]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	✓
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	✓
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	✓
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	✓
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	✓

In [35]:

```

1 ### Independent and Dependint features
2 ## Independent features= 13
3 X=dataset.iloc[:, :-1]      # lock all the columns from zero index to last except last t
4 X.head()

```

Out[35]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LS
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	5
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	5
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5

In [36]:

```

1 y= dataset.iloc[:, -1]   # only Last row
2 y.head()

```

Out[36]:

```

0    24.0
1    21.6
2    34.7
3    33.4
4    36.2
Name: Price, dtype: float64

```

**Note**

- Dependent feature is series
- Independent feature is dataframe

**Training and testing**

- Scaling will do after train and test of data set
- when we train the model then the model dont have any information about test data
- so train test split of the data is important
- [https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.model\\_selection.train\\_test\\_split.html](https://scikit-learn.org/stable/modules/generated/sklearn.model_selection.train_test_split.html))

In [37]:

```
1 from sklearn.model_selection import train_test_split
```

In [38]:

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=
```

- X\_train output is y\_train

- `X_test` output is `y_test`
- `X` is independent feature
- `y` is dependent feature
- `test_size` decides how much % of data goes to train data 33% of total data
- `random_state=42` my data is same with your if you select 42
- Size of `X_train` = `y_train` same for test

In [39]:

```
1 y_train.shape
```

Out[39]:

```
(339,)
```

In [40]:

```
1 X_train.shape
```

Out[40]:

```
(339, 13)
```

In [41]:

```
1 X_test.shape
```

Out[41]:

```
(167, 13)
```

In [42]:

```
1 y_test.shape
```

Out[42]:

```
(167,)
```

In [43]:

```
1 y_train
```

Out[43]:

```
478    14.6
26     16.6
7      27.1
492    20.1
108    19.8
...
106    19.5
270    21.1
348    24.5
435    13.4
102    18.6
Name: Price, Length: 339, dtype: float64
```

In [44]:

```
1 X_train
```

Out[44]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B
478	10.23300	0.0	18.10	0.0	0.614	6.185	96.7	2.1705	24.0	666.0	20.2	379.70
26	0.67191	0.0	8.14	0.0	0.538	5.813	90.3	4.6820	4.0	307.0	21.0	376.88
7	0.14455	12.5	7.87	0.0	0.524	6.172	96.1	5.9505	5.0	311.0	15.2	396.90
492	0.11132	0.0	27.74	0.0	0.609	5.983	83.5	2.1099	4.0	711.0	20.1	396.90
108	0.12802	0.0	8.56	0.0	0.520	6.474	97.1	2.4329	5.0	384.0	20.9	395.24
...	...	...	...	...	...	...	...	...	...	...	...	...
106	0.17120	0.0	8.56	0.0	0.520	5.836	91.9	2.2110	5.0	384.0	20.9	395.67
270	0.29916	20.0	6.96	0.0	0.464	5.856	42.1	4.4290	3.0	223.0	18.6	388.65
348	0.01501	80.0	2.01	0.0	0.435	6.635	29.7	8.3440	4.0	280.0	17.0	390.94
435	11.16040	0.0	18.10	0.0	0.740	6.629	94.6	2.1247	24.0	666.0	20.2	109.85
102	0.22876	0.0	8.56	0.0	0.520	6.405	85.4	2.7147	5.0	384.0	20.9	70.80

339 rows × 13 columns

## Feature Scaling or Standardization

- it use to reach globel minima (Gradent Decent)

In [45]:

```
1 from sklearn.preprocessing import StandardScaler
```

In [46]:

```
1 scaler= StandardScaler()
2 scaler      # StandardScaler() object is created
```

Out[46]:

StandardScaler()

In [47]:

```
1 ##### apply scaler on X train data
2 scaler.fit_transform(X_train)
```

Out[47]:

```
array([[ 0.89624872, -0.51060139,  0.98278223, ...,  0.86442095,
         0.24040357,  0.77155612],
       [-0.34895881, -0.51060139, -0.44867555, ...,  1.22118698,
         0.20852839,  0.32248963],
       [-0.41764058,  0.03413008, -0.48748013, ..., -1.36536677,
         0.43481957,  0.92775316],
       ...,
       [-0.43451148,  2.97567999, -1.32968321, ..., -0.56264319,
         0.36745216, -0.90756208],
       [ 1.01703049, -0.51060139,  0.98278223, ...,  0.86442095,
        -2.80977992,  1.50233514],
       [-0.40667333, -0.51060139, -0.38831288, ...,  1.17659123,
        -3.25117205, -0.26046005]])
```

In [48]:

```
1 ## after transfor store the data into variable called X_train or any new variable
2 X_train=scaler.fit_transform(X_train)
3 X_train
```

Out[48]:

```
array([[ 0.89624872, -0.51060139,  0.98278223, ...,  0.86442095,
         0.24040357,  0.77155612],
       [-0.34895881, -0.51060139, -0.44867555, ...,  1.22118698,
         0.20852839,  0.32248963],
       [-0.41764058,  0.03413008, -0.48748013, ..., -1.36536677,
         0.43481957,  0.92775316],
       ...,
       [-0.43451148,  2.97567999, -1.32968321, ..., -0.56264319,
         0.36745216, -0.90756208],
       [ 1.01703049, -0.51060139,  0.98278223, ...,  0.86442095,
        -2.80977992,  1.50233514],
       [-0.40667333, -0.51060139, -0.38831288, ...,  1.17659123,
        -3.25117205, -0.26046005]])
```

In [49]:

```
1 X_test=scaler.transform(X_test)
2 X_test
```

Out[49]:

```
array([[ -0.42451319, -0.51060139, -1.03649306, ..., -0.74102621,
         0.41899501, -0.48220406],
       [ -0.42911576,  1.2325393 , -0.6973123 , ..., -0.29506866,
         0.43481957, -1.25063772],
       [ -0.42269508, -0.51060139,  2.36824941, ...,  0.8198252 ,
         0.35807046,  0.77713459],
       ...,
       [ -0.33727525,  0.36096896, -1.04799071, ..., -2.34647337,
         0.38395492, -0.28556314],
       [ -0.30591027, -0.51060139, -0.44867555, ...,  1.22118698,
         0.2463943 , -0.07218683],
       [ -0.36872487,  0.36096896, -1.04799071, ..., -2.34647337,
         0.32133488, -0.91871901]])
```

### Important

- for test data only transform function is used not fit\_transform to avoid data leakage.
- model should not need to know which techniques are used for training.

## Model Training

### 1.Linear Regression Model

#### *Multiple regression used for?*

- making a prediction or forecasting
- [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html?highlight=linearregression#sklearn.linear\\_model.LinearRegression](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html?highlight=linearregression#sklearn.linear_model.LinearRegression) ([https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.LinearRegression.html?highlight=linearregression#sklearn.linear\\_model.LinearRegression](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LinearRegression.html?highlight=linearregression#sklearn.linear_model.LinearRegression))
- We can skip the transform and we can directly apply model training.
- in model training there is (normalize=True, default=False) if its True then transformation is done

In [50]:

```
1 from sklearn.linear_model import LinearRegression
```

In [51]:

```
1 reg = LinearRegression()  
2 reg
```

Out[51]:

LinearRegression()

- fit means training the data or to get the parameters. Data does not change
- fit\_transform means to set the data and change the data

In [52]:

```
1 reg.fit(X_train,y_train)  # y_train - dependent feature and X_train- independent feature
```

Out[52]:

LinearRegression()

In [53]:

```
1 ## print the coefficients and the intercept  
2 print(reg.coef_)
```

```
[-0.98858032  0.86793276  0.40502822  0.86183791 -1.90009974  2.80813518  
 -0.35866856 -3.04553498  2.03276074 -1.36400909 -2.0825356   1.04125684  
 -3.92628626]
```

In [54]:

```
1 ### there are 13 coefficients as features are 13 like crime, ZN, INDUS
```

In [55]:

```
1 print(reg.intercept_)
```

22.970796460176988

In [56]:

```
1 # if all the features are zero then the price of the house is 22.97  
2 # if one unit of increase in price then crime is decrease
```

## Prediction for test Data

In [57]:

```
1 reg_predict = reg.predict(X_test)
```



In [58]:

1 reg\_predict

Out[58]:

```
array([28.53469469, 36.6187006 , 15.63751079, 25.5014496 , 18.7096734 ,
       23.16471591, 17.31011035, 14.07736367, 23.01064388, 20.54223482,
       24.91632351, 18.41098052, -6.52079687, 21.83372604, 19.14903064,
       26.0587322 , 20.30232625,  5.74943567, 40.33137811, 17.45791446,
       27.47486665, 30.2170757 , 10.80555625, 23.87721728, 17.99492211,
       16.02608791, 23.268288  , 14.36825207, 22.38116971, 19.3092068 ,
       22.17284576, 25.05925441, 25.13780726, 18.46730198, 16.60405712,
       17.46564046, 30.71367733, 20.05106788, 23.9897768 , 24.94322408,
       13.97945355, 31.64706967, 42.48057206, 17.70042814, 26.92507869,
       17.15897719, 13.68918087, 26.14924245, 20.2782306 , 29.99003492,
       21.21260347, 34.03649185, 15.41837553, 25.95781061, 39.13897274,
       22.96118424, 18.80310558, 33.07865362, 24.74384155, 12.83640958,
       22.41963398, 30.64804979, 31.59567111, 16.34088197, 20.9504304 ,
       16.70145875, 20.23215646, 26.1437865 , 31.12160889, 11.89762768,
       20.45432404, 27.48356359, 10.89034224, 16.77707214, 24.02593714,
        5.44691807, 21.35152331, 41.27267175, 18.13447647,  9.8012101 ,
       21.24024342, 13.02644969, 21.80198374,  9.48201752, 22.99183857,
       31.90465631, 18.95594718, 25.48515032, 29.49687019, 20.07282539,
       25.5616062 ,  5.59584382, 20.18410904, 15.08773299, 14.34562117,
       20.85155407, 24.80149389, -0.19785401, 13.57649004, 15.64401679,
       22.03765773, 24.70314482, 10.86409112, 19.60231067, 23.73429161,
       12.08082177, 18.40997903, 25.4366158 , 20.76506636, 24.68588237,
        7.4995836 , 18.93015665, 21.70801764, 27.14350579, 31.93765208,
       15.19483586, 34.01357428, 12.85763091, 21.06646184, 28.58470042,
       15.77437534, 24.77512495,  3.64655689, 23.91169589, 25.82292925,
       23.03339677, 25.35158335, 33.05655447, 20.65930467, 38.18917361,
       14.04714297, 25.26034469, 17.6138723 , 20.60883766,  9.8525544 ,
       21.06756951, 22.20145587, 32.2920276 , 31.57638342, 15.29265938,
       16.7100235 , 29.10550932, 25.17762329, 16.88159225,  6.32621877,
       26.70210263, 23.3525851 , 17.24168182, 13.22815696, 39.49907507,
       16.53528575, 18.14635902, 25.06620426, 23.70640231, 22.20167772,
       21.22272327, 16.89825921, 23.15518273, 28.69699805,  6.65526482,
       23.98399958, 17.21004545, 21.0574427 , 25.01734597, 27.65461859,
       20.70205823, 40.38214871])
```

## Assumptions for Linear Regression

- assumptions are used to check model is working good or not

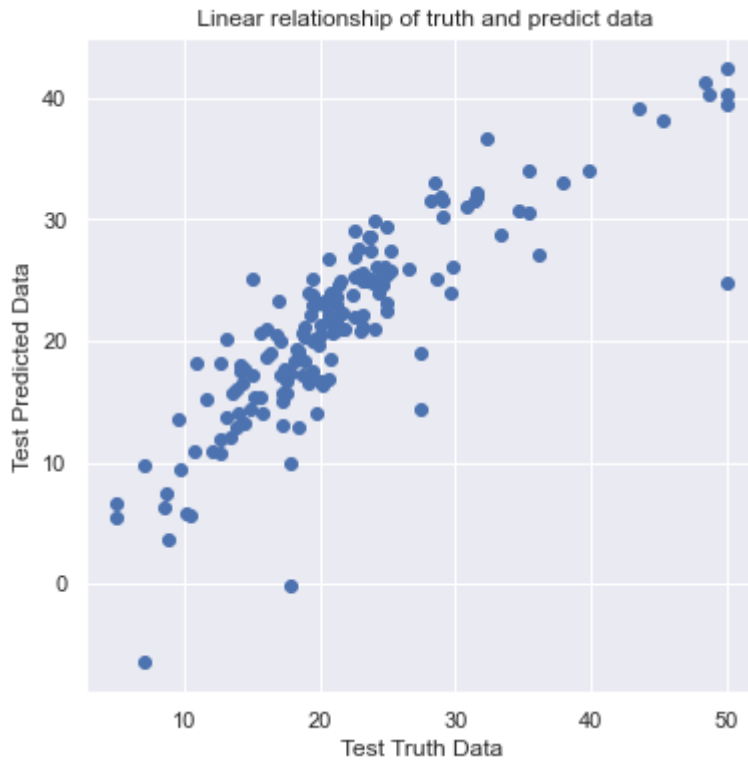
1 **### 1. First Assumption**

In [59]:

```
1
2 ## check scatter plot for y_test and predict data
3 plt.scatter(y_test,reg_predict)
4 plt.xlabel("Test Truth Data")
5 plt.ylabel("Test Predicted Data")
6 plt.title(" Linear relationship of truth and predict data")
```

Out[59]:

Text(0.5, 1.0, ' Linear relationship of truth and predict data')



- It shows linear relationship
- so model is good

## 2.Second Assumption

In [60]:

```
1 ### residuals means error
2 residuals= y_test- reg_predict
3 residuals
```

Out[60]:

```
173    -4.934695
274    -4.218701
491    -2.037511
72     -2.701450
452    -2.609673
```

```
...
```

```
110     0.642557
321    -1.917346
265    -4.854619
29     0.297942
262     8.417851
```


Name: Price, Length: 167, dtype: float64

In [61]:

```
1 sns.displot(residuals,kind="kde")
```

Out[61]:

<seaborn.axisgrid.FacetGrid at 0x206a059beb0>



### Observation

- the error plot is gaussian distribution (normal distribution) with negative skewed (right skewed) due to little outliers


## 3.Third Assumption

In [62]:

```
1 ### Scatter plot with preciction and residual  
2 ## Uniform distribution  
3 plt.scatter(reg_predict,residuals)  
4
```

Out[62]:

&lt;matplotlib.collections.PathCollection at 0x206a35eeb20&gt;



In [99]:

```
1 ## This above distribution is uniform distribution  
2 # doesn't have any define shape  
3
```

## 4.Fourth Assumption

In [64]:

```
1 ## performance Metric  
2
```

In [65]:

```
1 from sklearn.metrics import mean_squared_error  
2 from sklearn.metrics import mean_absolute_error
```

In [66]:

```

1 print("mean_squared_error : ", mean_squared_error(y_test,reg_predict))
2 print("mean_absolute_error:", mean_absolute_error(y_test,reg_predict))
3 print("Root mean square error : ",np.sqrt(mean_squared_error(y_test,reg_predict)))

```

```

mean_squared_error : 20.72402343733975
mean_absolute_error: 3.148255754816832
Root mean square error : 4.552364598463061

```

## R square and Adjusted R square

### R square

In [67]:

```

1 from sklearn.metrics import r2_score
2 score=r2_score(y_test,reg_predict)
3 print("R square Value is: ", score)

```

```
R square Value is: 0.7261570836552477
```

### Adjusted R2

- Formula
  - $\text{Adjusted R}^2 = 1 - [(1 - R^2) * (n - 1) / (n - k - 1)]$

In [68]:

```

1 ## Adjusted R square using formula
2 #display adjusted R-squared
3 Adjusted_R = 1 - (1-score)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
4 print("Adjusted R square is: ", Adjusted_R)

```

```
Adjusted R square is: 0.702889384880857
```

```

1 ##### Note :
2 * Adjusted R square is less than R square

```

## 2.Ridge Regression

- Ref : [https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Ridge.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html) ([https://scikit-learn.org/stable/modules/generated/sklearn.linear\\_model.Ridge.html](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.Ridge.html))
- `sklearn.linear_model.Ridge(alpha=1.0, *, fit_intercept=True, normalize='deprecated', copy_X=True, max_iter=None, tol=0.001, solver='auto', positive=False, random_state=None)`

In [71]:

```
1 from sklearn.linear_model import Ridge
2 ridge=Ridge()
3 ridge
```

Out[71]:

Ridge()

In [72]:

```
1 ridge.fit(X_train,y_train)
```

Out[72]:

Ridge()

In [129]:

```
1 ## print the coefficients and the intercept
2 print("***Ridge.coef are-- ",ridge.coef_)
3 print("***Ridge.intercept :",ridge.intercept_)
```

```
***Ridge.coef are-- [-0.97541551  0.84608896  0.37564928  0.86738391 -1.860
77739  2.81535042
-0.36108635 -3.00177053  1.95063015 -1.29462251 -2.06972563  1.03867858
-3.91121554]
***Ridge.intercept : 22.970796460176988
```

In [103]:

```
1 # prediction for test data
2 ridge_predict=ridge.predict(X_test)
```

```
1 ### Assumptions for Ridge Regression
```

## 1.First Assumption


- 1.Test truth data and Test predicted data should follow linear relationship.
- 2.If we get linear relationship then this indicates that the model is good.

In [76]:

```
1 plt.scatter(y_test,ridge_predict)
2 plt.xlabel("Test Truth Data")
3 plt.ylabel("Test Predicted Data")
4 plt.title(" Linear relationship of truth and predict data")
```

Out[76]:

Text(0.5, 1.0, ' Linear relationship of truth and predict data')



## 2.Second Assumption

- 1. Residuals should follow normal distribution.
- 2. If residuals follow normal distribution, it indicates that the model is good model.

In [110]:

```
1 ### residuals means error
2 residuals1= y_test- ridge_predict
3 residuals1.head()
```

Out[110]:


```
173    -4.907420
274    -4.166232
491    -2.143852
72     -2.691076
452    -2.604498
Name: Price, dtype: float64
```

In [108]:

```
1 sns.displot(residuals1,kind="kde")
```

Out[108]:

&lt;seaborn.axisgrid.FacetGrid at 0x206a55e3370&gt;



### 3.Third Assumption

- 1. Residuals vs Predictions should follow a uniform distribution.
- 2. If Residuals vs Predictions follow uniform distribution, it indicates this is good model.




In [107]:

```
1 ## Uniform distribution  
2 ### Scatter plot with preciction and residual  
3 plt.scatter(ridge_predict,residuals1)  
4 plt.xlabel('Predicted values')  
5 plt.ylabel('Residuals')
```

Out[107]:

Text(0, 0.5, 'Residuals')



## 4.Fourth Assumption

In [118]:

```

1 print("mean_squared_error : ", format(round(mean_squared_error(y_test,ridge_predict),4))
2 print("mean_absolute_error:", format(round(mean_absolute_error(y_test,ridge_predict),4))
3 print("Root mean square error : ",format(round(np.sqrt(mean_squared_error(y_test,ridge_

```

```

mean_squared_error : 20.7524
mean_absolute_error: 3.146
Root mean square error : 4.5555

```

## R square

In [121]:

```

1 from sklearn.metrics import r2_score
2 score1=r2_score(y_test,ridge_predict)
3 print("R square Value for ridge regression is: ", format(round(score1*100,4)))

```

```
R square Value for ridge regression is: 72.5782
```

## Adjusted R square

In [126]:

```

1 #adjusted R-squared
2 Adjusted_R_ridge = 1 - (1-score1)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
3 print("Adjusted R square for ridge regression is: ", format(round(Adjusted_R_ridge*100,

```

```
Adjusted R square for ridge regression is: 70.2482
```

```

1 ## 3.Lasso Regression
2 * REF : https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.Lasso.html?highlight=lasso
3 * lass sklearn.linear_model.Lasso(alpha=1.0, *, fit_intercept=True,
    normalize='deprecated', precompute=False, copy_X=True, max_iter=1000, tol=0.0001,
    warm_start=False, positive=False, random_state=None, selection='cyclic'

```

In [160]:

```

1 from sklearn.linear_model import Lasso
2 lasso = Lasso()
3 lasso
4 lasso.fit(X_train,y_train)
5 ## print the coefficients and the intercept
6 print("***Lasso.coef are-- ",lasso.coef_)
7 print("***Lasso.intercept :",lasso.intercept_)
8 print("_____")
9 # prediction for test data
10 lasso_predict=lasso.predict(X_test)
11 print("***lasso_predict",lasso_predict)
12 print("_____")
13 ### Assumptions of Lasso regression
14 ### 1.First Assumption ---- Linear relationship
15 plt.scatter(y_test, lasso_predict)
16 plt.xlabel('Test Truth data')
17 plt.ylabel('Test predicted data')
18
19 ## 2.second Assumption ---- Residual distribution
20 residuals_lasso=y_test-lasso_predict
21 print("***residuals_lasso.head()",residuals_lasso.head())
22 print("_____")
23
24 sns.displot(residuals_lasso, kind='kde')
25
26

```

```

***Lasso.coef are-- [-0.          0.          -0.          0.27140271 -0.
2.62932147
-0.          -0.          -0.          -0.          -1.21106809  0.29872625
-3.81788375]
***Lasso.intercept : 22.970796460176988

```

```

***lasso_predict [26.08015466 30.7480057 17.78164882 25.25224684 19.2838727
4 22.81161765
18.31125182 14.6359243 21.41277818 20.44276659 20.7857368 21.00978479
1.29101416 22.48591111 20.4207989 24.73115299 18.16643043 6.95747132
35.82658816 18.45664358 25.66618031 26.77096265 13.79601995 24.00317031
18.83677575 15.53225538 22.93567982 18.81410882 19.96419904 19.71394554
19.9929271 25.48086778 25.07506471 19.62299031 15.87164442 20.47826644
30.90020658 21.73740698 21.69357896 24.78795141 14.48946282 27.49872616
36.28097645 19.68302782 25.54695918 17.26691093 16.01035524 25.87512519
19.3705841 29.52965183 23.10173719 31.37342903 17.55332715 25.82107048
34.98857199 22.91267519 19.3967501 29.34678421 24.65125376 16.72971658
25.42537393 30.6751849 28.90511192 18.42571639 27.56426639 14.62706882
20.02272756 25.60745002 28.32959623 15.91971307 20.36020491 26.04012236
13.70562148 23.19186499 23.2538407 9.14791655 21.08680468 35.13203126
18.20120981 12.40579126 23.03574753 11.70030485 24.10234373 10.23869501
22.24788446 28.20852115 20.77401763 26.01572261 25.97666619 20.77471688
24.05595237 9.79658092 21.55718522 20.96232324 14.58941397 22.29462592
23.04513053 2.87810564 18.26028545 17.31405284 21.55660947 24.48282506
11.46772233 21.88129799 25.04349207 14.07796126 19.97841644 26.61705358
23.30429098 27.32736035 12.59741065 19.28050072 24.94727892 24.22470232
29.72519314 19.11634391 31.14895846 16.43050137 20.50890111 27.69026978
19.80307948 26.66386801 15.01321139 23.31466084 26.15446018 23.80801526
27.15999771 30.37432077 22.93935948 34.91159865 11.97264266 26.45153342
20.25377754 19.96681079 12.21677635 21.57200937 23.11587937 31.05309711

```

```
29.72484228 18.03669387 19.12012649 28.8679226 23.41788443 14.36679948
10.91433849 23.78530314 23.58885901 18.48704182 15.72569049 36.3867166
19.38880373 19.56932184 27.03174387 22.95041998 22.07807906 23.22702436
17.41528186 24.66493926 30.31735718 13.9998893 22.25446895 19.75004517
20.8350658 25.47389672 24.13577633 23.02944605 36.90324597]
```

---

```
***residuals_lasso.head() 173 -2.480155
```

```
274 1.651994
```

```
491 -4.181649
```

```
72 -2.452247
```


```
452 -3.183873
```


```
Name: Price, dtype: float64
```

---

Out[160]:

<seaborn.axisgrid.FacetGrid at 0x206a9268970>





In [171]:

```

1  ## 3. Third Assumption-----Uniform distribution
2  plt.scatter(lasso_predict, residuals_lasso)
3  plt.xlabel('Predicted values')
4  plt.ylabel('Residuals')
5
6  ### 4.Fourth Assumption
7  print("mean_squared_error : ", format(round(mean_squared_error(y_test,lasso_predict),4))
8  print("mean_absolute_error:", format(round(mean_absolute_error(y_test,lasso_predict),4))
9  print("Root mean square error : ",format(round(np.sqrt(mean_squared_error(y_test,lasso_
10 print("_____
11
12 ## R square
13 score2=r2_score(y_test,lasso_predict)
14 print("R square Value for lasso regression is: ", format(round(score2*100,4)))
15 print("_____
16
17 #adjusted R-squared
18 Adjusted_R_lasso = 1 - (1-score2)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
19 print("Adjusted R square for lasso regression is: ", format(round(Adjusted_R_lasso*100,
20

```

mean\_squared\_error : 26.1664  
mean\_absolute\_error: 3.6464  
Root mean square error : 5.1153

---

R square Value for lasso regression is: 65.4243

---




---

Adjusted R square for lasso regression is: 62.4865

---





## 4.Elastic-net Regression

In [170]:

```

1 from sklearn.linear_model import ElasticNet
2 Elesticnet = ElasticNet()
3 Elesticnet
4 Elesticnet.fit(X_train,y_train)
5 ## print the coefficients and the intercept
6 print("***ElasticNet.coef are-- ",Elesticnet.coef_)
7 print("***ElasticNet.intercept :",Elesticnet.intercept_)
8 print("
9 # prediction for test data
10 Elesticnet_predict=Elesticnet.predict(X_test)
11 print("***Elesticnet_predict",Elesticnet_predict)
12 print("
13 ### Assumptions of Lasso regression
14 ### 1.First Assumption ---- Linear relationship
15 plt.scatter(y_test, Elesticnet_predict)
16 plt.xlabel('Test Truth data')
17 plt.ylabel('Test predicted data')
18
19 ## 2.second Assumption ---- Residual distribution
20 residuals_Elesticnet=y_test-Elesticnet_predict
21 print("***residuals_Elesticnet.head()",residuals_Elesticnet.head())
22 print("
23
24 sns.displot(residuals_Elesticnet, kind='kde')
25

```

```

***ElasticNet.coef are-- [-0.36520114  0.          -0.14336748  0.63145824
-0.25193148  2.34999448
-0.          -0.          -0.          -0.25649969 -1.23951556  0.56384945
-2.56053213]
***ElasticNet.intercept : 22.970796460176988

```

```

***Elesticnet_predict [26.04802695 31.11448131 18.09845158 24.74715491 19.
13029713 23.07195028
19.8492127 16.42921582 20.98280883 21.03040905 23.59247585 22.4067143
2.50342106 22.86968897 21.05836477 23.53088819 19.32942155 9.24659633
34.51755093 18.33111982 25.39963891 26.53220506 16.04212388 23.68595117
18.22309609 15.9070075 22.91791506 17.40135861 22.80881602 20.34960072
21.28107265 25.0664737 23.29041734 18.52289666 16.68946719 20.17099878
29.78000437 22.08911412 24.00624402 24.52109601 16.51539744 27.25142517
34.8940966 20.75229792 25.54944362 17.27877681 17.51067948 25.422475
19.45141801 28.72445431 23.85816391 30.64335445 19.05778782 25.10137208
33.43673587 21.9368327 19.10068361 28.38705767 24.91075492 18.68821158
25.41735754 29.96236233 27.77368373 18.66077461 26.83456776 18.72984267
19.66634919 25.37569386 27.64862833 15.09326887 21.6230625 24.42218348
14.00935207 22.80340884 23.31057037 10.15388121 21.41270277 33.98445117
18.23197774 13.83630127 23.23840504 13.33915878 24.55297788 11.80396525
22.6039322 29.04777925 19.93864988 25.40796591 25.4253774 20.79626634
24.35937771 9.98031645 21.1883148 21.7010251 13.64158366 21.70329066
21.48780024 4.89921219 16.60651403 16.48691364 22.52662793 24.23810699
12.67234464 21.73288769 24.94869622 14.02785155 20.34560161 25.81816846
23.27665603 26.98968083 12.4461064 18.53919342 24.57036836 24.5991159
28.78917289 17.35695925 30.55890904 17.49669771 20.81635985 27.26668735
20.67056474 26.10145269 11.29182557 23.22360335 25.76790464 23.6220014
26.68354582 29.53505955 23.09099441 33.69060315 14.86234562 26.00700282
20.72314768 21.04261631 14.58583453 19.80159048 23.16541552 30.31837307
29.06727633 19.21664987 19.96431131 28.26947089 24.04394758 18.60022435

```




```
10.97392042 23.98834146 24.56611841 18.41001674 17.5657926 34.56534242
18.12003269 19.01050052 26.55906015 23.10161055 22.69122071 22.66142713
17.29393495 22.68213782 29.22756893 12.14001806 23.24525465 20.77586134
21.27148575 25.02037796 24.44501814 23.10652629 35.30973171]
```

```
***residuals_Elasticnet.head() 173 -2.448027
274 1.285519
491 -4.498452
72 -1.947155
452 -3.030297
Name: Price, dtype: float64
```

Out[170]:

<seaborn.axisgrid.FacetGrid at 0x206a943d250>





In [168]:

```

1  ## 3. Third Assumption-----Uniform distribution
2  plt.scatter(Elasticnet_predict, residuals_Elasticnet)
3  plt.xlabel('Predicted values')
4  plt.ylabel('Residuals')
5
6  ### 4.Fourth Assumption
7  print("mean_squared_error : ", format(round(mean_squared_error(y_test,Elasticnet_predict),4)))
8  print("mean_absolute_error:", format(round(mean_absolute_error(y_test,Elasticnet_predict),4)))
9  print("Root mean square error : ",format(round(np.sqrt(mean_squared_error(y_test,Elasticnet_predict),4)))
10 print("_____")
11
12 ## R square
13 score3=r2_score(y_test,Elasticnet_predict)
14 print("R square Value for Elasticnet regression is: ", format(round(score2*100,4)))
15 print("_____")
16
17 #adjusted R-squared
18 Adjusted_R_Elastic = 1 - (1-score2)*(len(y_test)-1)/(len(y_test)-X_test.shape[1]-1)
19 print("Adjusted R square for Elastnet regression is: ", format(round(Adjusted_R_Elastic,4)))
20

```

```

mean_squared_error :  27.1402
mean_absolute_error: 3.6277
Root mean square error :  5.2096

```

---

```

R square Value for Elasticnet regression is:  65.4243

```

---




---

```

Adjusted R square for Elastnet regression is:  62.4865

```

---



In [ ]:

```
1
```

END

In [169]:

```
1 ## Hope You Like The Presentation Skills
```